

Testing the Sattelzeit as a Network Event

A method note on the **Koselleck Machine**

August 4, 2026

Contents

1	Context	2
2	Method	2
2.1	Corpus	5
3	The interactive explorer	5
4	Current limitations	7

1 Context

The historian Reinhart Koselleck argued that the period 1770 – 1830 (the **Sattelzeit**, German for “saddle period”) was not simply an era in which many words happened to change meaning at the same time. He claimed it was a moment when the whole vocabulary of politics and society reorganized **as a system** – concepts shifted together, in relation to one another, rather than drifting independently one at a time. This has always been argued in prose, from close reading of texts. It has never been tested computationally.

Ryan Heuser’s *Computing Koselleck* took the first step towards such a test: he trained a separate word-embedding model for each historical period and measured how far each individual word moves from one period to the next. He found that this movement is unusually large around 1770 – 1830, which supports the idea that something real happened in that window. But his method looks at one word at a time. It cannot say whether words moved **together**, as a system reorganizing around a shared structure, or whether each word’s shift is its own independent event that simply happens to cluster in time.

This project adds that second, missing test. Instead of asking how far a single word moves, it asks whether the **groupings** among words – the way the whole vocabulary is organized into clusters of related meaning – change unusually around the same decades. If Koselleck is right, we should see the cluster structure itself reorganize at the Sattelzeit, not just individual words drifting off on their own.

2 Method

The method reuses the same per-period training that Heuser’s approach uses, so the two studies remain directly comparable, and adds a network layer on top that a word-by-word approach cannot produce.

1. **Split the corpus into time windows.** The corpus is divided into even 20-year windows, currently 1500 to 1900. The windows are uniform across the whole timeline, not just around the Sattelzeit: if the method measured more finely right where reorganization was expected and more coarsely everywhere else, it would be biased toward finding exactly what it was looking for. Heuser’s own results, in fact, show a second spike of word-level change around 1905 – 1925, unrelated to the Sattelzeit – a pattern that a Sattelzeit-only sampling scheme could never have picked up. The 1500 – 1900 span reflects the corpus assembled for this project so far, not a limit of the method itself – the same pipeline applies unchanged to any other span or window size, set entirely in configuration, given a period-dated corpus to feed it.

2. **Train a language model per period.** For each time window, a model (**word2vec**) is trained independently, learning a numeric representation of every word from how it is actually used in that period's text: words that appear in similar contexts end up with similar representations. Training a separate model per period, rather than one model updated continuously, is what makes it possible to later measure both individual word drift and whole-system reorganization: a word's meaning only makes sense relative to the other words trained alongside it, in that same period.
3. **Build a word-similarity network.** From those representations, a network is built for each period: every word is a node, linked to its 15 closest neighbours by **cosine similarity** (a standard measure of how alike two words' representations are). A fixed number of nearest neighbours is used instead of a similarity threshold, because at this vocabulary scale a fixed threshold produces networks that are almost fully connected – tens of millions of edges per period, unusable and with no meaningful structure to cluster. Purely grammatical words (articles, prepositions, pronouns) and single-letter tokens are excluded as network nodes: the interest is in words that carry conceptual content, not the grammatical scaffolding of the language (those grammatical words are still used during training itself, since they help place every other word correctly).
4. **Detect communities.** A **community detection** algorithm (**Leiden**) is run on each network, grouping words that are densely connected to one another. In practice, these groups function as semantic fields: sets of words that, in that period, are used in similar ways to each other more than to the rest of the vocabulary.

This differs from familiar clustering methods such as *k*-means in a basic way. *k*-means fixes a number of groups in advance and assigns each item to whichever group's average position it sits closest to. Leiden does neither: it does not fix the number of groups beforehand, and it never compares a word to a group's "average" – it only ever asks a relational question about a word's actual connections. For any word, and any group it is already directly connected to, Leiden checks whether moving that word into the group would make the network as a whole denser *within* groups and sparser *between* them than plain chance would predict, given how connected that word and that group already are overall. If the move helps, it is made; if not, the word stays put or a different move is tried. Repeating this, word by word, until no single move can improve things any further, is the entire algorithm.

The number this process is trying to improve at every step has a name, **modularity** (written *Q*). It compares how connected the words placed inside the same group actually are against how connected they would be expected to be if the same set of connections had been handed out at random instead. A grouping that captures real structure scores well above zero; a grouping no better than chance scores around zero; a grouping that forces genuinely unrelated words together scores *below* zero,

actively worse than not grouping at all. Leiden's search is nothing more than hill-climbing on this one number.

5. **Repeat detection at several levels of detail.** Community detection does not have one single “correct” answer: a parameter called **resolution** controls how many groups appear and how large they are, by changing how large a word's own connections have to be, relative to chance, before joining a group is judged worthwhile. For that reason every network is processed at seven different resolution settings, and any claim about reorganization has to hold up across most of those settings, not just one chosen after the fact. Critically, no single setting among the seven is ever treated as “the” result and the rest discarded – the sweep's entire purpose is to check that the same conclusion holds broadly across settings, not to pick whichever one happens to agree with the hypothesis. This is a hard rule of the project: no finding is reported if it only holds at a single, convenient resolution. (One setting, resolution 1.0 – plain textbook modularity, with no bias toward larger or smaller groups – is used separately, only to choose a single picture to display in the interactive explorer described below. That is a display choice with no bearing on which finding gets reported as real, and it is not an arbitrary one: looking across the already-computed sweep, modularity – how well a grouping captures real structure, against random chance – sits on a broad plateau from resolution 0.5 to 1.5 and peaks, only marginally, at 1.0, while the migration fraction rises and period-to-period stability falls fairly steadily as resolution increases beyond a low setting. A higher display resolution would therefore mechanically inflate the very reorganization measurement this project is testing, which is reason enough to keep the display fixed at the standard default rather than a value chosen after seeing the result.)
6. **Compare consecutive periods.** For each pair of consecutive periods, the comparison is restricted to words that appear in both (each period has its own vocabulary). The best possible match between one period's groups and the next period's groups is found first – group numbers do not mean the same thing from one period to the next, so the correspondence has to be established before anything can be compared – and then the fraction of shared words that still ended up in a different group, despite that best match, is measured. That fraction, how much of the shared vocabulary “migrated” to a new group, is what this project calls the **migration fraction** – its central measurement of how much reorganization happened between two periods.
7. **Test whether reorganization concentrates at the Sattelzeit.** The hypothesis predicts that the migration fraction should be unusually high specifically for the transitions that fall inside or close out the 1770 – 1830 window, compared to every other transition across the full 1500 – 1900 timeline.

8. **Repeat the whole measurement per region.** The corpus is not evenly British across the timeline, and the imbalance is worst exactly where it matters most: in 1780 – 1800 the American material (Evans-TCP, 1681 documents) outweighs the British (ECCO-TCP, 910 documents), because ECCO’s public transcriptions thin out sharply at the end of the eighteenth century. A change in the cluster structure across that boundary could therefore reflect a change in *who was writing* rather than a change in the language. To separate the two, every step above – training, network construction, community detection, the resolution sweep, and the migration fraction – is also run on each region’s documents alone, in addition to the combined corpus. If the reorganization is real, it should show up inside the British-only and the American-only timelines as well, not only in the mixture of them; if it shows up only in the combined run, it is an artifact of corpus composition rather than semantic history. Nothing here hardcodes which regions exist: a document’s region is simply the top-level folder its source was filed under, so the same check extends to any other split of the corpus.
9. **Dictionary cross-check (not yet implemented).** As a second line of evidence, independent of the network entirely, the plan is to take the words that change group right at the pivot and check whether the OED or other period dictionaries independently record a dated sense change at the same moment. This layer is not yet designed: the corpus mixes British and American sources, and matching each source against the wrong regional dictionary would invalidate the comparison, so the choice of dictionary per source has to be resolved first.

2.1 Corpus

The primary source is the **Text Creation Partnership (TCP)**: curated, manually corrected transcriptions with no OCR noise, covering 1500 – 1800 (EEBO-TCP, ECCO-TCP, Evans-TCP). Since the Sattelzeit extends to 1830 and TCP stops at 1800, the plan is to supplement 1800 – 1900 with Project Gutenberg. HistWords (pre-trained historical embeddings) is used only as an independent sanity check, not as primary evidence. Each source is filed under a region (currently `british` for EEBO-TCP and ECCO-TCP, `american` for Evans-TCP), which is what makes the per-region rerun described above possible at all.

3 The interactive explorer

Alongside the pipeline, a small web tool – the **Koselleck Machine** – makes the per-period networks directly explorable, without reading raw numbers. It has three views of the same underlying data.

Timeline. The primary view. Type a word and see, in one continuous strip, which

community it belonged to in every period it appears in – the community’s plain-English label and its lane (below) as a secondary colour cue, an explicit gap where the word is simply absent from that period’s vocabulary, and a marked step wherever the word’s community actually changed from the period before, versus stayed the same. Earlier drafts of this tool led with the graph explorer below, but historian collaborators who reviewed it found a force-directed diagram hard to read at a glance; what they consistently asked for instead was exactly this – which domain a word belonged to, over time, read directly rather than inferred from a picture. Clicking any period in the strip opens that word’s full neighbourhood there, reusing the word search view below.

Word search. The same underlying data, presented as a plain table instead of a diagram: type a word and a period, and get back its closest neighbours by similarity, each neighbour’s community, and whether the searched word’s own community changed since the previous period. Useful when the actual numbers matter more than the picture, and reused directly by the timeline’s own per-period drill-in above.

Graph explorer. A period slider drives a force-directed network diagram: nodes are words, their size reflects how connected they are, and their color marks which community they belong to in that period. A search box focuses the view on one word and its neighbourhood; moving the slider one period at a time shows that neighbourhood evolving, and a change in a word’s color between periods marks a real change of community, not just a coincidence of the diagram’s layout. A side panel explains the method in plain terms and shows, for the currently selected resolution, how the migration measurement holds up across all seven resolution settings – the same robustness check described above, made visible rather than just claimed. Kept for browsing the raw network structure – and because building it first is what surfaced this whole labeling and lane apparatus – but no longer the primary way to read the tool, per the reception above.

Region toggle. Where the per-region runs described above have been built, both views also offer a region switch – combined, British only, American only – so the same word in the same period can be compared across the sub-corpora directly, and a shift can be checked against the possibility that it only exists in the mixture. The switch lists only the regions that were actually built, read off the data itself rather than assumed.

Every community shown in any view carries a short, plain-English label (for example “Government & Law”, “Moral Character & Virtue”). These labels were generated by having a language model read each community’s most representative words and summarize the shared theme in a few words – a convenience for a human reader, not a finding in themselves. A label is not re-read independently every period, however: whenever a community’s best correspondence to the period before it (the same correspondence step the migration fraction is built on) can be established, its label is inherited unchanged from that predecessor, and a fresh read only happens for a community with no such predecessor – the first period a region’s data begins, or the side of a genuine

reorganization that a word moved into. In the combined corpus, roughly seven in eight community labels are inherited this way rather than independently generated for a given pipeline run; the per-region runs behave similarly. This keeps a community's displayed name consistent with whether the tool has actually flagged it as having moved or stayed, rather than the two being computed separately and occasionally disagreeing, which an early version of the tool did. Roughly 60 – 70% of communities are honestly flagged, by the same process, as **“mixed”**: the corpus is multilingual (English, Latin, Law French, Welsh, Scots, and more) and early modern, and a community detected by this method sometimes groups words by shared grammar – verb conjugations, comparative forms, spelling variants – rather than by shared topic. Every view surfaces this caveat directly next to the labels, rather than presenting every label as equally reliable.

Each community is additionally sorted into one **lane**: a small, fixed list of about a dozen broad domains (government and law, religion, medicine, trade, and so on), plus one explicit “structural / uncertain” lane for exactly the “mixed” case above. The lane list itself was authored once by hand, from a read-through of every non-mixed label already produced, then fixed – a language model assigns each community to one of these existing lanes, it never invents a new one per run. That fixed-in-advance property is what lets a lane stay a meaningful, comparable signal across independent pipeline runs even though, as the next section explains, the underlying community structure and the raw label text do not reproduce identically run to run. In the tools above, a community's plain-English label and number stay the primary signal; its lane is a secondary colour cue layered on top, not a replacement for reading the actual label.

4 Current limitations

Three limitations shape what can be claimed right now, and all three are visible in the tool itself rather than hidden.

Corpus coverage. The corpus does not yet cover 1800 – 1900 with the same density as earlier periods, since the Gutenberg supplement described above has not been sourced yet. The periods with the least surviving text are, not coincidentally, the ones closing the Sattelzeit window, so any comparison across that boundary currently rests on a thinner sample than the rest of the timeline. That gap skews composition and not only volume: with ECCO thin after 1780, the 1780 – 1800 window holds more American than British documents, which is why the per-region rerun exists. It bounds the problem rather than solving it – each regional sub-corpus is thinner still than the combined one, so per-region results carry wider uncertainty than the headline numbers.

Label reliability. A community's label is inherited, not re-read, for as long as its best correspondence to the previous period holds – only a minority of labels (roughly one in ten to one in six, depending on region) are an independent read for a given pipeline

run. That independent read is still a single, unvalidated pass by a language model, not a checked ground truth – useful for orientation, not for citation. Inheritance carries its own, different cost: on the corpus’s longest-running community lineages, a label can go up to fifteen periods (close to three hundred years) without being independently re-read, and the community’s actual word composition can drift well past what the label still says by then – a staleness problem, not a wrong-read one, and not yet addressed.

Reproducibility. Rerunning the full pipeline from scratch on the same corpus does not reproduce an identical result, and this is worth stating plainly rather than glossing over. Community detection itself is properly seeded and deterministic given a fixed network – the gap sits one step earlier, in training: language-model training here uses multiple parallel threads for speed, and that kind of parallel training is not bit-for-bit repeatable even with a fixed random seed, because the exact order in which threads update shared values depends on how the operating system happens to schedule them that run. Small differences in the trained word vectors cascade into a similarity network with slightly different edges, which cascades into a structurally similar but not identical community partition. In practice this shows up as, for instance, one full rerun finding 304 communities in the combined corpus where an earlier run found 308 – a real difference, not a bug, and not one worth trading away the threefold training speed of parallel execution to chase, since it does not bear on the actual claim under test. That claim is whether reorganization concentrates at the Sattelzeit, tested by the resolution sweep described above, holding across settings within one run – not whether a specific community keeps the same id, or even the same member words down to the last one, across two independently retrained copies of the pipeline. The one thing designed to stay meaningfully comparable across independent reruns is the small, fixed lane list described above, precisely because it is authored once by hand and never reinvented per run.