

THE KOSELLECK MACHINE

The Koselleck Machine

A project overview: what this instrument measures, who it is for, what it can and cannot support yet, and how to run or extend it.

REPOSITORY	<code>koselleck-networks</code>
DOCUMENT	<code>docs/overview.tex</code> → <code>docs/overview.pdf</code>
COMPANION	<code>docs/method.pdf</code> (the method in full)
STATUS	research instrument, local only, not yet deployed
COMPILED	August 7, 2026

Contents

1 What this is 2

2 Who it is for, and what it is for right now 2

3 What is in the repository 3

4 Where it stands 4

5 Running it 6

 5.1 Setup 6

 5.2 Running the pipeline 7

 5.3 Running the web tool locally 8

 5.4 Bringing your own corpus 8

6 People 9

7 Where to read more 9

1 What this is

The Koselleck Machine is a research instrument built to test one historical claim, and a small web tool for reading the result.

The claim is Reinhart Koselleck's. He argued that the decades between roughly 1770 and 1830, which he called the **Sattelzeit** (German for “saddle period”), were not merely a stretch of time in which many words happened to change meaning at once, but a moment when the vocabulary of politics and society reorganized **as a system**: concepts moved in relation to one another rather than drifting independently. That has always been argued in prose, from close reading. Ryan Heuser's *Computing Koselleck* was the first computational test of it: he trained a separate word-embedding model per historical period and measured how far each word moved from one period to the next, and found that movement is unusually large around 1770-1830. But a one-word-at-a-time measurement cannot distinguish a system reorganizing from a pile of unrelated shifts that happen to coincide in time. This project adds the missing test. It builds a word-similarity network for each period, runs community detection on it, and measures whether the **grouping structure itself** reorganizes at the Sattelzeit, rather than only measuring individual words.

The full argument, the algorithms, and the exact measurements are set out in the companion document, `docs/method.pdf`. This overview does not repeat them. Its job is orientation: what the thing is, who it is for, what state it is in, and how to run it.

2 Who it is for, and what it is for right now

Two readers are in scope.

A historian or intellectual historian who wants to interrogate the result rather than take a summary statistic on faith. For that reader, the web tool is the product: type a word, see which semantic community it belonged to in every twenty-year window from 1500 onward, see explicitly where it moved and where it stayed, and drill into any period to read the actual nearest neighbours behind that judgement.

A computational or digital-humanities reader who wants to check the method, disagree with it, or run it on different text. For that reader, the product is the pipeline: seven standalone scripts, a single versioned configuration file, and a corpus contract small enough that substituting your own sources is a folder drop or one reader function (Section 5.4).

What it is not, yet. This is a research instrument, not a finished, citation-ready dataset. The numbers it produces are the object under test, not a published result; the community labels are a reading aid, not a validated taxonomy; and a full rerun does not reproduce an identical partition. Section 4 states each of those plainly. Treat anything read out of the tool as a lead to check, not as evidence to cite.

3 What is in the repository

src/ is the pipeline, run in order: `parse_tcp.py` (raw sources to normalised per-document records), `bucket_periods.py` (records to 20-year text windows), `embeddings.py` (one word2vec model per window), `network.py` (each word linked to its 15 nearest neighbours by cosine similarity), `community.py` (Leiden community detection at seven resolution settings), and `metrics.py` (how much the grouping changed between consecutive periods: migration fraction, NMI, adjusted Rand). Alongside them sit `pipeline_config.py`, the shared configuration loader, and the two optional labeling scripts described below.

webapp/ is the Koselleck Machine itself, a Flask application with a landing page and three views onto the same pre-built networks: `/timeline` (the primary view: one word's community across every period, in a single strip), `/graph` (a D3 force-directed explorer for one word's neighbourhood in one period, plus a sampled full-network mode), and `/search` (a plain lookup table: nearest neighbours, community, and whether the word's community changed since the previous period). Where the pipeline was also run per region, every page gains a region toggle, listing only the regions that were actually built.

docs/ holds this overview and the method note.

labels/ holds a small, citable snapshot of the current community labels, one CSV and one compiled JSON per region, copied there by `label_communities.py` `publish` so they travel with the code instead of living only in the gitignored data directory.

config.yml holds the settings that are shared and versioned: the period slices, and the word2vec and Leiden hyperparameters. Period boundaries and model parameters are configuration here, not code.

What the repository deliberately does *not* ship is the corpus or the trained models. That is a size decision, not a rights one: the raw Text Creation Partnership zips and the derived per-period networks run to many gigabytes, and the sources are public domain and fetchable directly. The contribution is the measurement, not the data.

4 Where it stands

Six limitations bound what can be claimed today. All six are stated in full in `docs/method.pdf` and surfaced inside the tool itself rather than hidden; they are summarised here without softening.

Corpus coverage. The corpus is the Text Creation Partnership (EEBO-TCP, ECCO-TCP, Evans-TCP, curated and manually corrected, no OCR noise, 1500-1800), supplemented for 1800-1900 by the British Library's digitised 19th-century books collection (public domain, OCR-derived, catalogue metadata attached). A Project Gutenberg supplement was considered instead and dropped: its metadata carries only its own digitisation date, not the book's original publication date, which the method needs to bucket documents by period at all. The supplement is filed under the `british` region, alongside EEBO/ECCO - not because American and British English were different languages for most of this span (they were not, and treating the split as a dialect claim would not be defensible), but because it is a different, independently curated archive, so the region rerun tests whether the reorganization signal survives across two provenance sources rather than one, and the combined corpus stays the headline result. A direct consequence: the `american` region has no data past 1800 at all, since Evans-TCP stops there and the supplement is British-only. ECCO's own thinness after 1780 also still means the 1780-1810 window leans American-heavy even with the supplement in place. Per-region results carry wider uncertainty than the combined ones, and the American arm specifically cannot speak to anything after the Sattelzeit's closing edge.

OCR artifacts. Unlike TCP, the British Library supplement is OCR-derived, and OCR text carries its own artifacts. One class - an original line-break hyphen left sitting in the text ("*utrum- que*" for *utrumque*) - was found and fixed. A second, distinct class was not: the OCR sometimes drops the line-break hyphen entirely and leaves a bare space instead ("*par ticulars*" for *particulars*), with no punctuation signal left to detect it by. This produces whole communities that are nothing but word fragments rather than a real topic - confirmed directly in the 1810-1830 period, at least fifteen such communities, some several hundred words large. The labeling pass handles this honestly (routed to "Structural / Uncertain," not mislabeled), but the underlying vocabulary and network metrics for 1810-1830, 1830-1850, and 1870-1890 are still diluted by it. A real fix, not yet built, most plausibly needs a dictionary check rather than a punctuation pattern.

A separate, unconfirmed TCP artifact. A third word-splitting pattern, distinct from the OCR issue above, turned up in a handful of TCP-covered periods (1510-1650) while raising the labeling resolution (next box) - TCP is manually keyed, not scanned, so it cannot be the same bug. Words are missing a “con-”/“com-” prefix (*mytted* for *committed*, *uersation* for *conversation*), most likely from how TCP’s own markup for a scribal abbreviation mark gets handled during extraction - not confirmed, not fixed, flagged rather than guessed at.

Labeling resolution, raised once already. The webapp displays and labels communities at one fixed resolution, separate from the sweep every quantitative claim is tested across. It was 1.0 until the British Library supplement pushed some periods’ vocabularies past 170,000 words and the largest community past 24,000 words - too large and mixed for any label to describe. Raised to 4.0 after a concrete check: the community holding “system” in the worst-affected period stayed an unstructured mix at 1.0 through 3.0, and only split into a real, nameable theme (civil administration and property law) at 4.0. The resolution sweep used for the actual claim was extended to match (up to 4.0). Expect this to need raising again as the corpus grows further - it scales with vocabulary size, it is not a fixed constant.

Label reliability. Every community carries a short plain-English name, and those names are generated by a language model reading the community’s most representative words. A label is inherited, not re-read, for as long as the same alignment that decides “moved”/“stayed” maps a community back to a predecessor; only a minority of labels (roughly one in ten to one in six, depending on region) are an independent read in a given run. That independent read is a single, unvalidated pass, useful for orientation and not for citation. Inheritance carries a different cost: on the longest-running lineages a label can go up to fifteen periods, close to three hundred years, without being re-read, while the community’s actual word composition drifts past what the label still says. Separately, roughly half of communities carry the “Structural / Uncertain” lane rather than a topical one (down from 70% before the resolution was raised, see below), because the corpus is multilingual (English, Latin, Law French, Welsh, Scots, and more) and early modern, and a detected community sometimes groups words by shared grammar or shared language rather than by shared topic. The tool shows that lane next to the label rather than presenting every label as equally reliable.

Reproducibility. Rerunning the pipeline from scratch on the same corpus does not reproduce an identical result. Community detection is properly seeded and deterministic given a fixed network; the gap sits one step earlier, in training, which uses multiple parallel threads and is therefore not bit-for-bit repeatable even with a fixed seed, because thread scheduling varies. Small differences in the vectors cascade into slightly different edges and then into a structurally similar but not identical partition: one full rerun found 304 communities in the combined corpus where an earlier run found 308. That is a real difference, not a bug, and it is not worth trading away the threefold speed of parallel training to chase, because the claim under test is whether reorganization concentrates at the Sattelzeit and holds across the seven resolution settings within a run, not whether a specific community keeps its id across two independently retrained copies. The one thing designed to stay comparable across reruns is the fixed lane list, because it is authored by hand once and never reinvented.

There is also a standing methodological rule worth repeating here, because it governs everything the project will ever report:

Hard rule. Community detection resolution changes the answer. Every network is therefore processed at seven resolution settings, and no finding is reported if it holds only at a single, convenient one. The sweep exists to test that a conclusion survives across settings, not to pick the setting that agrees with the hypothesis.

Finally, on availability: the Koselleck Machine currently runs locally only. A public hosted deployment is intended but has not been started, deliberately kept out of scope until the local application is solid. The Flask app runs as-is on any host that can run Python, but it is not a static site and cannot be served from GitHub Pages, since every response is computed server-side from the network data.

5 Running it

5.1 Setup

```
python -m venv .venv
.venv/Scripts/activate (or source .venv/bin/activate on macOS or Linux)
pip install -r requirements.txt
```

The pipeline expects a `data_root` folder containing `corpus/`, `processed/`, `embeddings/`,

networks/, and communities/. Point at your own copy either by creating a gitignored `config.local.yml` with a `data_root:` entry, or by setting the `DATA_ROOT` environment variable, which takes priority over both config files. Only `corpus/` is filled in by hand; everything else is produced by the pipeline. The exact folder names come from `config.yml`'s `paths` block, which is authoritative if anything is renamed.

The corpus itself is fetched separately. All three TCP components used here are public domain, and the bulk P4 XML zip shards are downloaded directly from the Text Creation Partnership. They are laid out under `corpus/tcp/regions/<region>/<source>/*.zip`, for example `british/eebo_phase1/`, `british/ecco/`, `american/evans/`. Both folder levels are free-form: `parse_tcp.py` reads the zips in place, without unpacking them, and discovers regions and sources by scanning that tree rather than from any hardcoded list. The top-level folder name becomes the region tag used throughout the pipeline and the webapp's region toggle. Only quality-checked "released" shards are read; TCP's "unedited" variants are skipped on purpose.

5.2 Running the pipeline

Each stage is a standalone script, run in order:

```
python src/parse_tcp.py
python src/bucket_periods.py
python src/embeddings.py
python src/network.py
python src/community.py
python src/metrics.py
```

That sequence is enough to reproduce every quantitative result. Community labels are a separate, optional stage on top of it, run as two more scripts:

```
python src/extract_community_words.py
python src/label_communities.py generate --region combined
python src/label_communities.py compile --region combined
python src/label_communities.py publish --region combined
```

`generate` produces a human-editable CSV, inheriting each community's label from its aligned predecessor wherever one exists and leaving only genuinely new communities blank for a person or an agent to fill in; `compile` turns the CSV into the JSON the webapp reads; and `publish` copies the CSV and JSON into the repository's `labels/` directory as a citable snapshot. `--region` also accepts a region name or `all`. The prompt the model is given lives in the repository as `src/prompts/community_labeling_v1.md`, and `docs/method.pdf` quotes and documents it in full.

5.3 Running the web tool locally

```
python webapp/app.py
```

Then open `http://127.0.0.1:5000`. Without real data the app still starts, but every period shows as a coverage gap. For a production host, the app is served with `gunicorn` and a `DATA_ROOT` environment variable pointing at wherever the network and community data lives on that machine; the data does not travel with the repository and has to be uploaded or fetched there separately.

5.4 Bringing your own corpus

The pipeline is built so that a reader can rerun the whole measurement on their own text rather than take these numbers on trust. There are exactly two cases, and only one of them touches code.

Case A: the material is already in TCP's P4 XML format (another TCP release, or a repackaged TCP subset). No code change at all. Drop the zip files at `<data_root>/corpus/tcp/regions/<a-region-name>/<a-source-name>/*.zip` and rerun the pipeline from `src/parse_tcp.py`. Nothing needs registering anywhere: both folder names are free-form, the region level becomes the region tag carried through the pipeline and the webapp, the source level is a diagnostic label kept in the manifest, and `parse_tcp.py` discovers both by scanning.

Case B: the material is in any other raw format (plain text, JSON, a different XML flavour, OCR dumps). One file changes: `src/parse_tcp.py`. Add a reader next to the existing `iter_xml_members` that walks the new format and, for each document, emits the same record the TCP path already writes to `processed/all_docs.jsonl`:

```
{"region": "...", "source": "...", "doc_id": "...", "year": 1782, "text": "...}"
```

plus the matching `region`, `source`, `doc_id`, `year`, `chars` row in `manifest.csv`. Documents with no extractable year or no text are dropped, as in the TCP path. That is the entire contract. Everything downstream, `bucket_periods.py` through `metrics.py` and the webapp, reads only those normalised records and needs no change; period boundaries and model hyperparameters are configuration in `config.yml`, not code. In both cases, a region name that did not exist before is picked up automatically by the region-split outputs and the webapp's region toggle; there is no list of valid regions to update.

To be explicit about what does not exist: there is no upload interface, no hosted ingestion service, and no API-key-driven pipeline. The mechanism is the ordinary one, clone the repository, put your corpus files in the right folder, rerun the scripts in order.

6 People

Ryan Heuser wrote *Computing Koselleck*, the word-level antecedent this project extends, and is the author of the base method.

Jamie McGarry pointed to the public TCP corpus and holds Cambridge institutional access to the fuller restricted EEBO, ECCO, and Evans archives beyond the public TCP subset, if a future window needs more text than public TCP alone provides.

Bernardo Villegas supervises the project and has given the running feedback on method and on the design of the tool.

7 Where to read more

`docs/method.pdf` is the full method note: the argument, the algorithms, the resolution sweep and why it is non-negotiable, the labeling prompt and workflow in full, and the limitations in their unabridged form. The repository's `README.md` carries the operational detail: corpus download links and licensing, the exact expected data layout, and deployment notes.